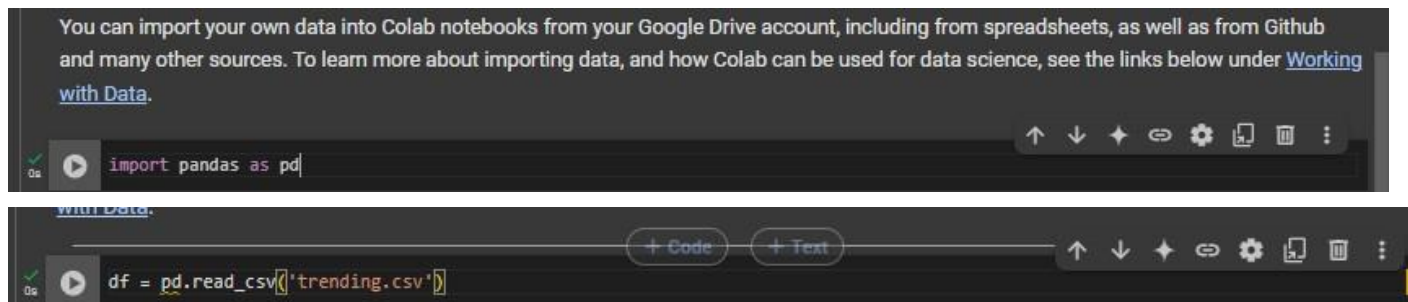


## DS-1 Lab Exp 1

**AIM:** Introduction to Data science and Data preparation using Pandas steps.

- Load data in Pandas.
- Description of the dataset.
- Drop columns that aren't useful.
- Drop rows with maximum missing values.
- Take care of missing data.
- Create dummy variables.
- Find out outliers (manually)
- standardization and normalization of columns

Step 1: Firstly import Pandas Library as pd and then Load data in Pandas using `pd.read_csv`.

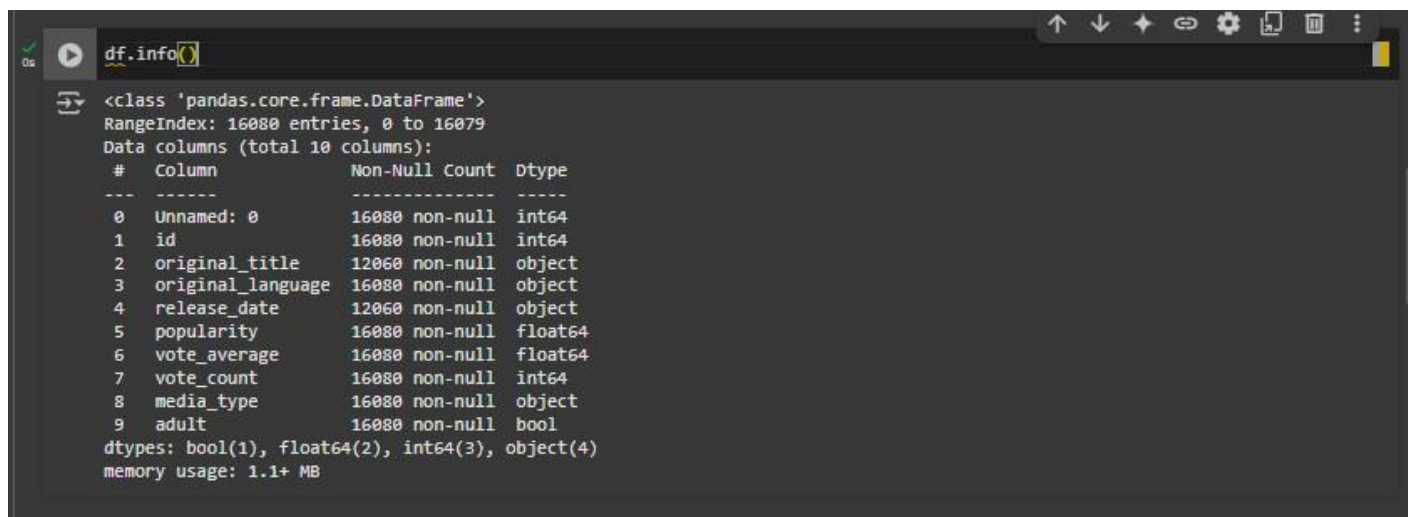


The screenshot shows two code cells in a Google Colab notebook. The first cell contains the code `import pandas as pd`. The second cell contains the code `df = pd.read_csv('trending.csv')`. Above the code cells, there is a message about importing data from Google Drive, GitHub, and other sources.

Step 2: Get Description of the Dataset by using following 2 commands

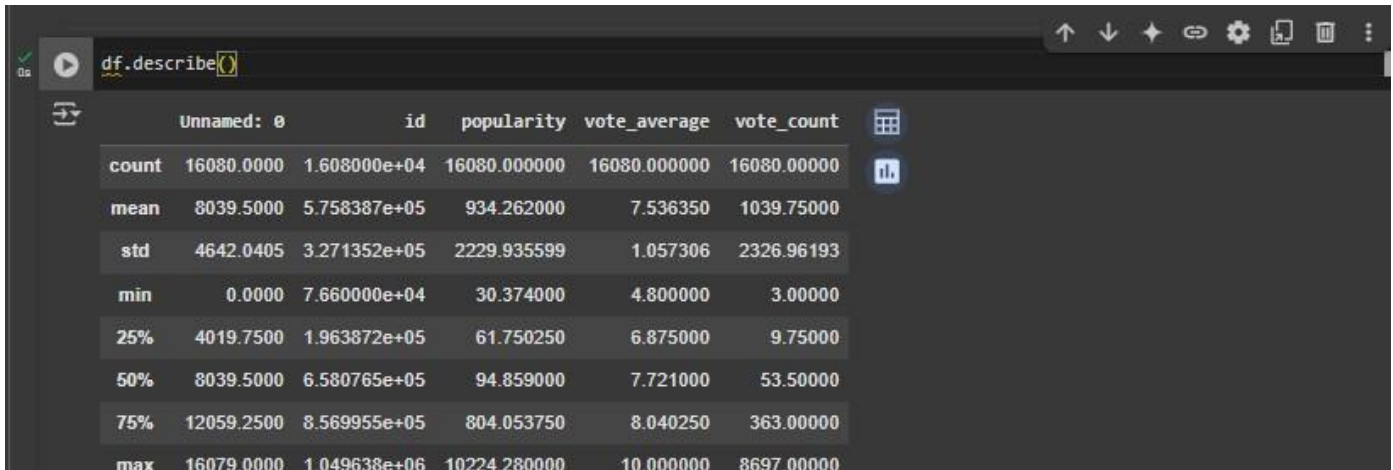
`df.info()` -> Get basic information about the dataset

`df.describe()` -> Summary statistics of the dataset



The screenshot shows a code cell in a Google Colab notebook with the command `df.info()`. The output of the command is displayed below the code cell. It shows the class of the DataFrame, the range of the index, the total number of columns, and a table of columns with their non-null counts and data types. At the bottom, it shows the dtypes and memory usage.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16080 entries, 0 to 16079
Data columns (total 10 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   Unnamed: 0          16080 non-null  int64
 1   id                  16080 non-null  int64
 2   original_title      12060 non-null  object
 3   original_language   16080 non-null  object
 4   release_date        12060 non-null  object
 5   popularity          16080 non-null  float64
 6   vote_average        16080 non-null  float64
 7   vote_count          16080 non-null  int64
 8   media_type          16080 non-null  object
 9   adult              16080 non-null  bool
dtypes: bool(1), float64(2), int64(3), object(4)
memory usage: 1.1+ MB
```



A screenshot of a Jupyter Notebook cell showing the output of the `df.describe()` command. The output is a summary statistics table for the DataFrame. The columns are: Unnamed: 0, id, popularity, vote\_average, and vote\_count. The rows represent different statistical measures: count, mean, std, min, 25%, 50%, 75%, and max.

|       | Unnamed: 0 | id           | popularity   | vote_average | vote_count   |
|-------|------------|--------------|--------------|--------------|--------------|
| count | 16080.0000 | 1.608000e+04 | 16080.000000 | 16080.000000 | 16080.000000 |
| mean  | 8039.5000  | 5.758387e+05 | 934.262000   | 7.536350     | 1039.75000   |
| std   | 4642.0405  | 3.271352e+05 | 2229.935599  | 1.057306     | 2326.96193   |
| min   | 0.0000     | 7.660000e+04 | 30.374000    | 4.800000     | 3.00000      |
| 25%   | 4019.7500  | 1.963872e+05 | 61.750250    | 6.875000     | 9.75000      |
| 50%   | 8039.5000  | 6.580765e+05 | 94.859000    | 7.721000     | 53.50000     |
| 75%   | 12059.2500 | 8.569955e+05 | 804.053750   | 8.040250     | 363.00000    |
| max   | 16079.0000 | 1.049638e+06 | 10224.280000 | 10.000000    | 8697.00000   |

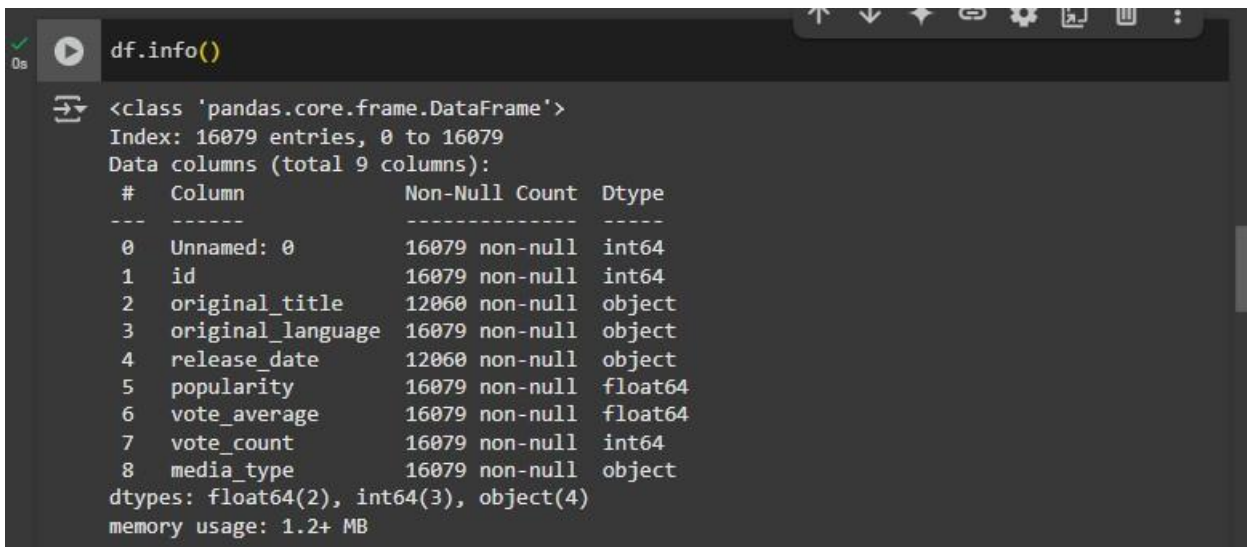
Step 3: Drop Columns that aren't useful. From Our Dataset we are dropping the "adult" column .



A screenshot of a Jupyter Notebook cell showing the code to drop the 'adult' column from the DataFrame. The code is: `cols = ['adult']` followed by `df = df.drop(cols,axis=1)`.

```
cols = ['adult']
df = df.drop(cols,axis=1)
```

We can see that it returned total 9 columns as it dropped the adult column



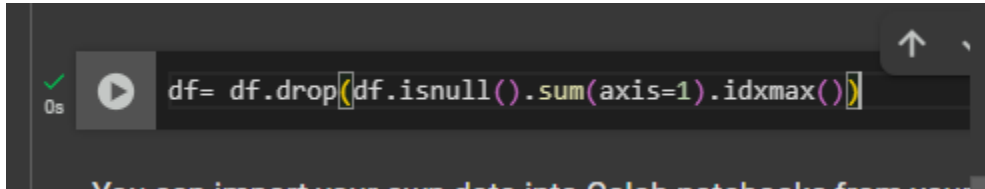
A screenshot of a Jupyter Notebook cell showing the output of the `df.info()` command. The output shows the DataFrame has 16079 entries and 9 columns. The columns are: Unnamed: 0, id, original\_title, original\_language, release\_date, popularity, vote\_average, vote\_count, and media\_type. The dtypes are: float64(2), int64(3), object(4). The memory usage is 1.2+ MB.

```
<class 'pandas.core.frame.DataFrame'>
Index: 16079 entries, 0 to 16079
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Unnamed: 0            16079 non-null  int64  
1   id                    16079 non-null  int64  
2   original_title        12060 non-null  object  
3   original_language     16079 non-null  object  
4   release_date          12060 non-null  object  
5   popularity            16079 non-null  float64 
6   vote_average          16079 non-null  float64 
7   vote_count            16079 non-null  int64  
8   media_type            16079 non-null  object  
dtypes: float64(2), int64(3), object(4)
memory usage: 1.2+ MB
```

Step 4: Drop row with maximum missing values.

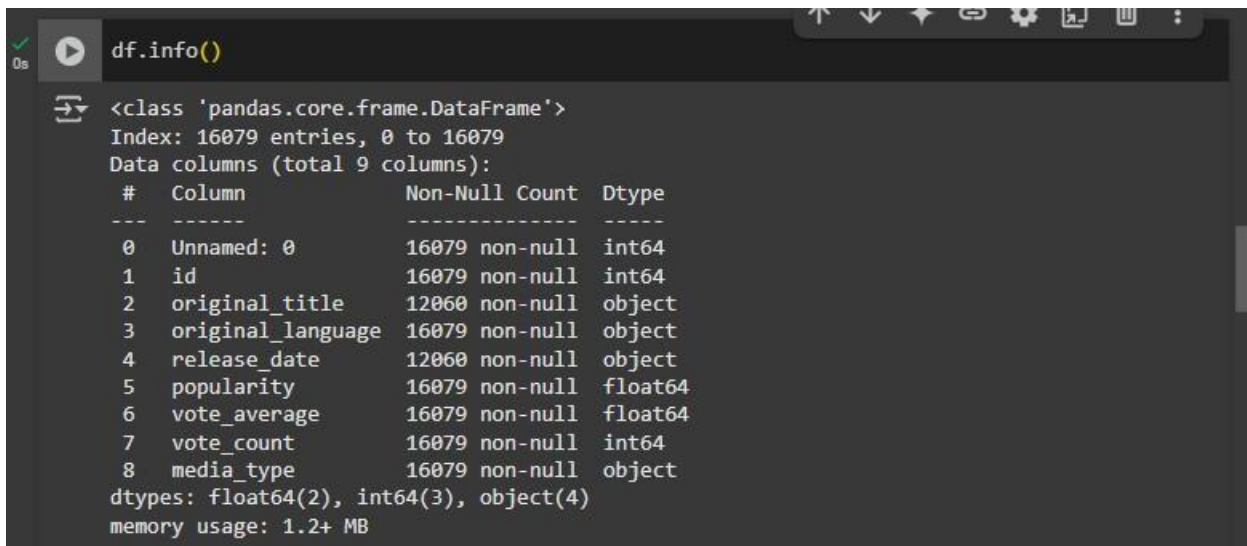
`df.isnull().sum(axis=1)` -> Computes the number of missing values (NaN) for each row.

`.idxmax()` -> Returns the index of row with max. no. of missing value



```
df= df.drop(df.isnull().sum(axis=1).idxmax())
```

We can see below that `df.info()` returns total 16079 entries, initially there were 16080 entries

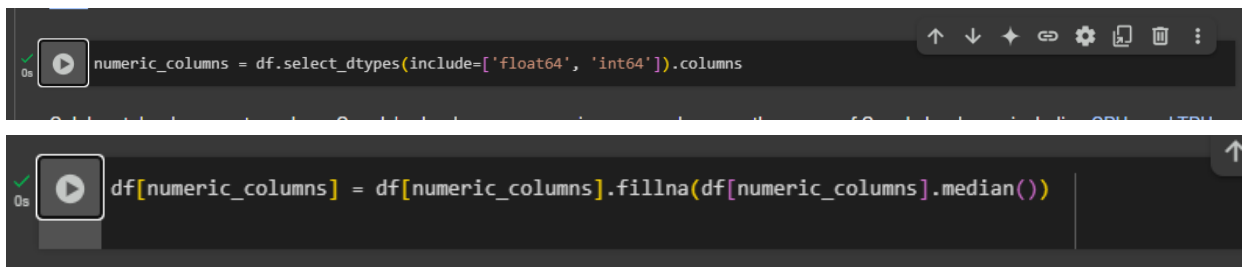


```
df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 16079 entries, 0 to 16079
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Unnamed: 0            16079 non-null  int64  
 1   id                    16079 non-null  int64  
 2   original_title        12060 non-null  object  
 3   original_language     16079 non-null  object  
 4   release_date          12060 non-null  object  
 5   popularity            16079 non-null  float64 
 6   vote_average          16079 non-null  float64 
 7   vote_count            16079 non-null  int64  
 8   media_type            16079 non-null  object  
dtypes: float64(2), int64(3), object(4)
memory usage: 1.2+ MB
```

Step 5: Taking care of missing data.

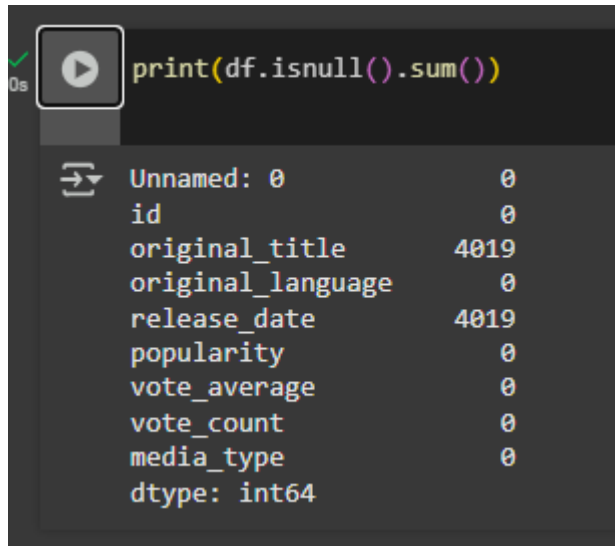
We can fill the empty numeric values with mode or median or mean. Below we had filled it with median. Firstly we had fetched the numeric values and then using `.fillna().median` we had filled it.



```
numeric_columns = df.select_dtypes(include=['float64', 'int64']).columns

df[numeric_columns] = df[numeric_columns].fillna(df[numeric_columns].median())
```

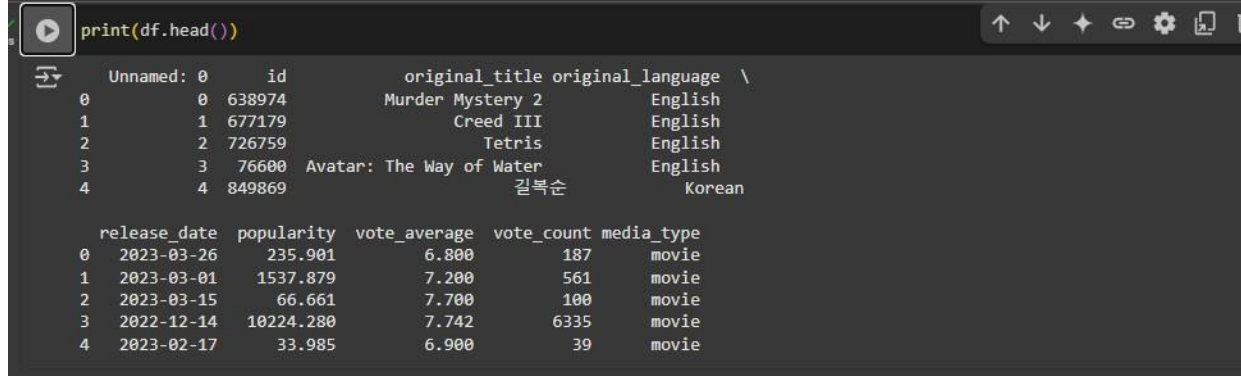
We can see that all the columns which had empty are filled. As they returned the sum 0



```
print(df.isnull().sum())
```

| Unnamed: 0        | 0     |
|-------------------|-------|
| id                | 0     |
| original_title    | 4019  |
| original_language | 0     |
| release_date      | 4019  |
| popularity        | 0     |
| vote_average      | 0     |
| vote_count        | 0     |
| media_type        | 0     |
| dtype:            | int64 |

df.head() returns starting 5 values



```
print(df.head())
```

|   | Unnamed: 0 | id     | original_title           | original_language | \ |
|---|------------|--------|--------------------------|-------------------|---|
| 0 | 0          | 638974 | Murder Mystery 2         | English           |   |
| 1 | 1          | 677179 | Creed III                | English           |   |
| 2 | 2          | 726759 | Tetris                   | English           |   |
| 3 | 3          | 76600  | Avatar: The Way of Water | English           |   |
| 4 | 4          | 849869 | 김복순                      | Korean            |   |

|   | release_date | popularity | vote_average | vote_count | media_type |
|---|--------------|------------|--------------|------------|------------|
| 0 | 2023-03-26   | 235.901    | 6.800        | 187        | movie      |
| 1 | 2023-03-01   | 1537.879   | 7.200        | 561        | movie      |
| 2 | 2023-03-15   | 66.661     | 7.700        | 100        | movie      |
| 3 | 2022-12-14   | 10224.280  | 7.742        | 6335       | movie      |
| 4 | 2023-02-17   | 33.985     | 6.900        | 39         | movie      |

```

0s print(df.head(20))
6      6      493529 Dungeons & Dragons: Honor Among Thieves
7      7      932430 Prom Pact
9      9      816904 Momias
10     10     514999 Murder Mystery
11     11     1049638 Rye Lane
12     12     739405 Operation Fortune: Ruse de Guerre
13     13     158876 NaN
14     14     921355 Assassin
15     15     117465 NaN
16     16     933419 Champions
17     17     208891 NaN
18     18     878375 On a Wing and a Prayer
19     19     82856 NaN
20     20     638974 Murder Mystery 2

original_language release_date popularity vote_average vote_count \
0      English      2023-03-26      235.901      6.800      187
1      English      2023-03-01     1537.879      7.200      561
2      English      2023-03-15      66.861      7.700      100

```

0s completed at 12:09 PM

Step 6: Create dummy variables. By using the below commands separate columns are created for each unique value in a column

```

0s df = pd.get_dummies(df)
+ Code + Text

```

```

0s print(df.head(20))
7      False      False      False
9      False      False      False
10     False      False      False
11     False      True      False
12     False      False      False
13     False      False      False
14     False      False      False
15     False      False      False
16     False      False      False
17     False      False      False
18     False      False      False
19     False      False      False
20     False      False      False

release_date_2023-03-23 release_date_2023-03-26 release_date_2023-03-30 \
0      False      True      False

```

We can understand the working here,

As we can see that we now it have returned 42 columns. But previously our data had 9 columns .

So this change is because of the dummy variables , it have created separate column for each unique value in a column

Below it shows original\_title\_Assassin, original\_language\_English.

```
df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 16079 entries, 0 to 16079
Data columns (total 42 columns):
#   Column                                                                 Non-Null Count  Dtype
---  -
0   Unnamed: 0                                                            16079 non-null  int64
1   id                                                                    16079 non-null  int64
2   popularity                                                            16079 non-null  float64
3   vote_average                                                          16079 non-null  float64
4   vote_count                                                            16079 non-null  int64
5   original_title_Assassin                                              16079 non-null  bool
6   original_title_Avatar: The Way of Water                            16079 non-null  bool
7   original_title_Champions                                             16079 non-null  bool
8   original_title_Creed III                                             16079 non-null  bool
9   original_title_Dungeons & Dragons: Honor Among Thieves             16079 non-null  bool
10  original_title_John Wick: Chapter 4                                  16079 non-null  bool
11  original_title_Momias                                                 16079 non-null  bool
12  original_title_Murder Mystery                                         16079 non-null  bool
13  original_title_Murder Mystery 2                                       16079 non-null  bool
14  original_title_On a Wing and a Prayer                                16079 non-null  bool
15  original_title_Operation Fortune: Ruse de Guerre                     16079 non-null  bool
16  original_title_Prom Pact                                              16079 non-null  bool
17  original_title_Rye Lane                                               16079 non-null  bool
18  original_title_Tetris                                                  16079 non-null  bool
19  original_title_김복순                                                 16079 non-null  bool
20  original_language_Chinese                                             16079 non-null  bool
21  original_language_English                                             16079 non-null  bool
```

✓ 0s completed at 12:21 PM



### Step 7: Create Outliers

They identify and handle unusual values in a dataset.

We are using Z-score to handle the data

```
{x} 0s ✓ from scipy import stats

# Select only numerical columns
numerical_df = df.select_dtypes(include=['float64', 'int64'])

# Remove constant or problematic columns
numerical_df = numerical_df.loc[:, numerical_df.nunique() > 1]
numerical_df = numerical_df.dropna(axis=1)

# Calculate Z-scores
z_scores = stats.zscore(numerical_df)

# Handle cases with NaN Z-scores
z_scores = pd.DataFrame(z_scores, columns=numerical_df.columns).fillna(0)

# Identify rows with Z-scores > 3 or < -3
outliers = (abs(z_scores) > 3).any(axis=1)

# Filter the outliers
outlier_rows = df[outliers]
print(outlier_rows)
```

|       | Unnamed: 0 | id    | popularity | vote_average | vote_count | \     |
|-------|------------|-------|------------|--------------|------------|-------|
| 3     | 3          | 76600 | 10224.280  | 7.742        | 6335       |       |
| 19    | 19         | 82856 | 1108.646   | 8.488        | 8697       |       |
| 23    | 23         | 76600 | 10224.280  | 7.742        | 6335       |       |
| 39    | 39         | 82856 | 1108.646   | 8.488        | 8697       |       |
| 43    | 43         | 76600 | 10224.280  | 7.742        | 6335       |       |
| ...   | ...        | ...   | ...        | ...          | ...        | ...   |
| 16039 | 16039      | 82856 | 1108.646   | 8.488        | 8697       |       |
| 16043 | 16043      | 76600 | 10224.280  | 7.742        | 6335       |       |
| 16059 | 16059      | 82856 | 1108.646   | 8.488        | 8697       |       |
| ..... | .....      | ..... | .....      | .....        | .....      | ..... |

### Step 8: Standardization and Normalization

Import StandardScaler and MinMaxScaler

```
✓ 0s [23] from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

Standardization (z-score scaling) transforms the data by subtracting the mean and dividing by the standard deviation for each feature.

```

# Select numerical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns

# Initialize the StandardScaler
scaler = StandardScaler()

# Standardize the numerical columns
df[numerical_columns] = scaler.fit_transform(df[numerical_columns])

# Check the results
print(df.head())

```

|   | Unnamed: 0 | id        | popularity | vote_average | vote_count | \ |
|---|------------|-----------|------------|--------------|------------|---|
| 0 | -1.732158  | 0.192916  | -0.313201  | -0.696417    | -0.366495  |   |
| 1 | -1.731943  | 0.309711  | 0.270665   | -0.318094    | -0.205769  |   |
| 2 | -1.731727  | 0.461279  | -0.389096  | 0.154808     | -0.403883  |   |
| 3 | -1.731512  | -1.526286 | 4.166043   | 0.194532     | 2.275593   |   |
| 4 | -1.731296  | 0.837632  | -0.403749  | -0.601836    | -0.430097  |   |

|   | original_title_Assassin | original_title_Avatar: The Way of Water | \ |
|---|-------------------------|---|---|
| 0 | False                   | False                                   |   |
| 1 | False                   | False                                   |   |
| 2 | False                   | False                                   |   |
| 3 | False                   | True                                    |   |
| 4 | False                   | False                                   |   |

Normalization scales numerical data to a fixed range, usually [0, 1]. Use MinMaxScaler for this process.

```

# Initialize the MinMaxScaler
scaler = MinMaxScaler()

# Normalize the numerical columns
df[numerical_columns] = scaler.fit_transform(df[numerical_columns])

# Check the results
print(df.head())

```

|   | Unnamed: 0 | id       | popularity | vote_average | vote_count | \ |
|---|------------|----------|------------|--------------|------------|---|
| 0 | 0.000000   | 0.577957 | 0.020162   | 0.384615     | 0.021164   |   |
| 1 | 0.000062   | 0.617220 | 0.147883   | 0.461538     | 0.064182   |   |
| 2 | 0.000124   | 0.668174 | 0.003560   | 0.557692     | 0.011157   |   |
| 3 | 0.000187   | 0.000000 | 1.000000   | 0.565769     | 0.728318   |   |
| 4 | 0.000249   | 0.794696 | 0.000354   | 0.403846     | 0.004141   |   |

|   | original_title_Assassin | original_title_Avatar: The Way of Water | \ |
|---|-------------------------|---|---|
| 0 | False                   | False                                   |   |
| 1 | False                   | False                                   |   |
| 2 | False                   | False                                   |   |
| 3 | False                   | True                                    |   |
| 4 | False                   | False                                   |   |



**Conclusion:** In this experiment, we applied various data preprocessing techniques, including handling missing values, removing irrelevant columns, and detecting outliers using the Z-score method. We then scaled the numerical data using standardization (Z-score method) and normalization (Min-Max scaling) to bring all features onto a uniform scale.

**Some Challenges we faced :**

1. Handling Missing Data: Identifying the appropriate method to handle missing values and replacing them with mean, median, or mode.
2. Scaling and Normalization: Deciding between standardization and normalization for different features can be tricky. Using incorrect scaling methods may distort the data and affect model accuracy.
3. Selection of Columns: Determining which columns are relevant for the model and dropping them is challenging.